




LUND
UNIVERSITY

Lecture 3 – Machine Learning Fundamentals



Outline – Lecture 3


- Artificial Neural Networks for Deep Learning
- Neuronal activation functions
- Backpropagation and Gradient descent
- Deep Belief Nets



Motivation



The "atom" of an Artificial Neural network



Output

Weighted sum of inputs

$$\hat{y} = \sigma \left(b + \sum_{i=1}^m x_i w_i \right)$$


Non-linear activation function

"Bias"


The Activation function, $\sigma(z)$

- **Q:** Why do we need an activation function? → **Nonlinearity** (allows modelling of nonlinear systems)


$$\hat{y} = \sigma \left(b + \sum_{i=1}^m x_i w_i \right)$$



$$\sigma(z) = \frac{1}{1 + e^{-z}}$$


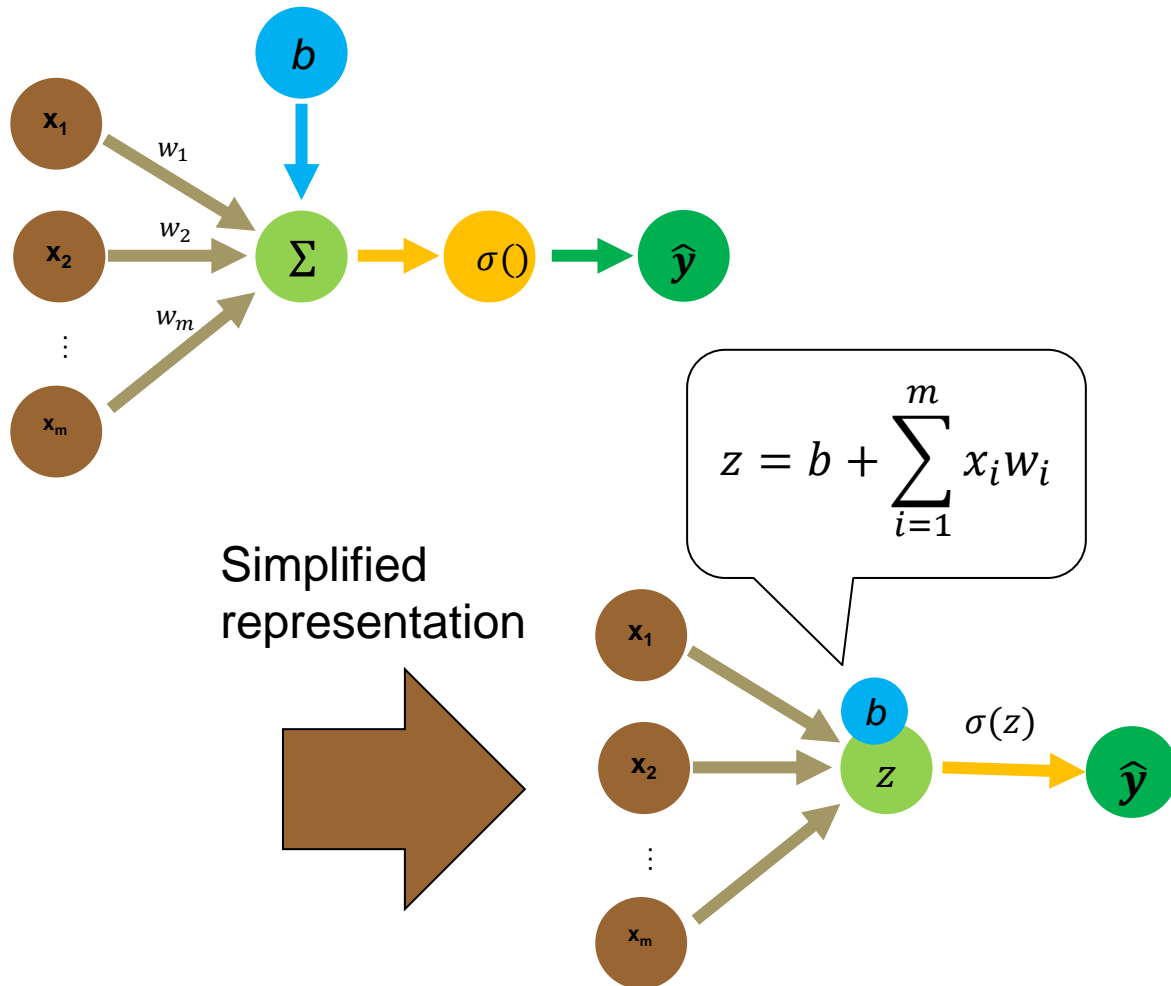


$$\sigma(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

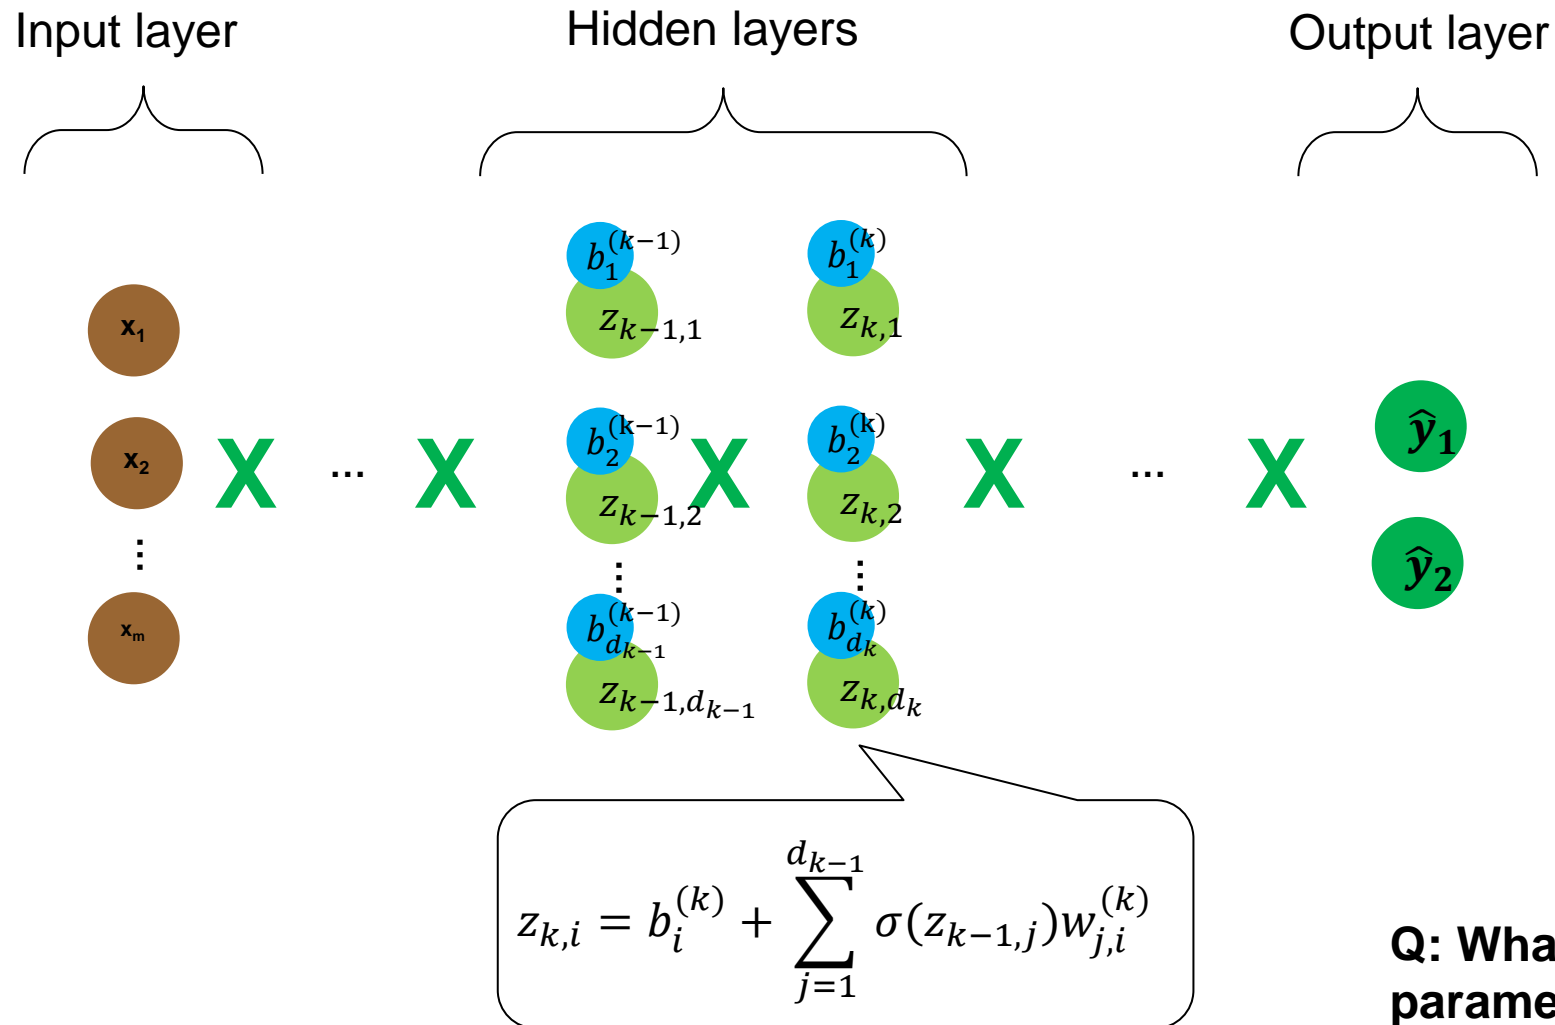


$$\sigma(z) = \max(0, z)$$

Single layer Neural network

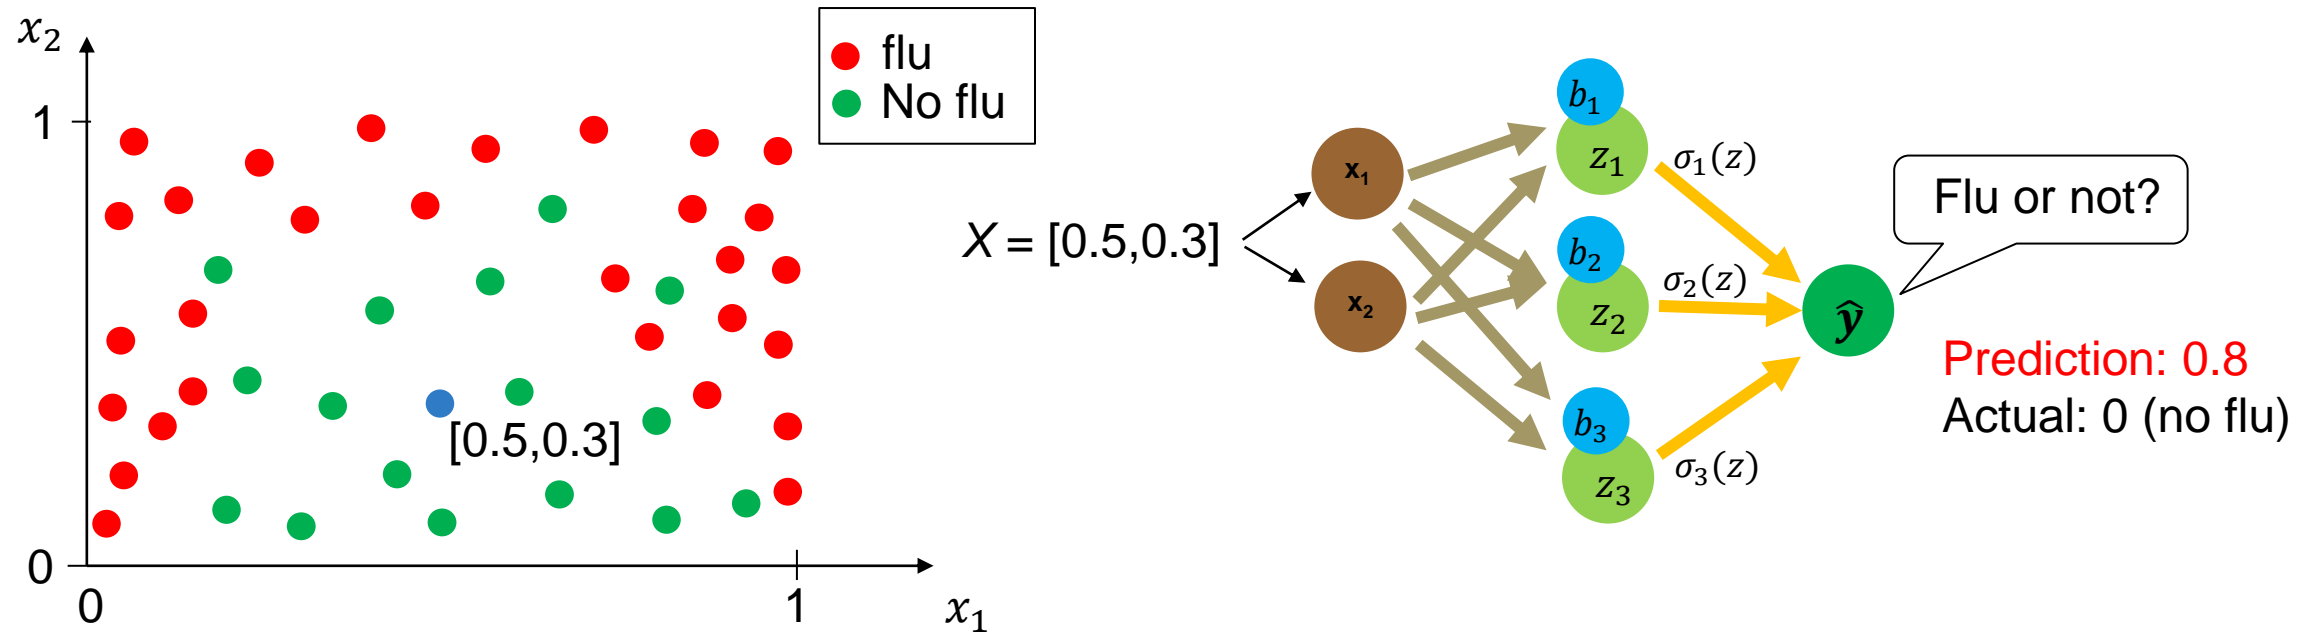


Deep Neural Network




Example use case

- Task: Network to predict likelihood of getting the flu
- x_1 : Time of the year (normalized) (i.e. 1st Jan = 0, 31st Dec = 1)
- x_2 : Time spent with other people (0 = never, 1 = 24/7)



Cost functions and training principle

- Cross-entropy cost function: $C(W, b) = -\frac{1}{n} \sum_{i=1}^n (t \ln \hat{y} + (1 - t) \ln(1 - \hat{y}))$ For **discrete** output
- Mean Squared Error Cost function: $C(W, b) = \frac{1}{n} \sum_{i=1}^n (t^i - f(X^i, W, b))^2$ For **continuous** output
- Training means finding the (W, b) that minimizes C for our dataset.



1. Initialize W and b
2. Calculate $\nabla C(W, b)$
3. Update W and b according to $W = W - \eta \nabla C$
4. Repeat 2-3 until convergence


Learning rate

Go in direction of negative gradient

So how to calculate ∇C ?


Matrix representation

$$W_l = \begin{bmatrix} w_{11} & \cdots & w_{1,n} \\ \vdots & \ddots & \vdots \\ w_{m,1} & \cdots & w_{m,n} \end{bmatrix}$$



$$A_l = \begin{bmatrix} a_1^l \\ \vdots \\ a_m^l \end{bmatrix}$$

$$B_l = \begin{bmatrix} b_1^l \\ \vdots \\ b_m^l \end{bmatrix}$$



The weighted input of the l^{th} layer:

$$Z_l = W_l A_{l-1} + B_l$$

The activation of the l^{th} layer is thus:

$$A_l = \sigma(Z_l)$$

Backpropagation algorithm

- A way to calculate ∇C

1. **Input X:** Set the activation A_1 at the input layer
2. **Feed-forward:** for each $l = 2, 3, \dots, L$ compute Z_l and A_l

3. **Output error:**

Compute $C(W_L, B_L)$ and $\frac{\partial C}{\partial Z_L} = \frac{\partial C}{\partial A_L} * \sigma'(Z_L)$ of the layer L

$$\rightarrow \frac{\partial C}{\partial W_L} = \frac{\partial C}{\partial Z_L} A_{L-1}$$

4. **Back-propagate:** For each layer $l < L$, calculate

$$\frac{\partial C}{\partial W_l} = \frac{\partial C}{\partial Z_l} \frac{\partial Z_l}{\partial W_l} = \underbrace{\left(W_{l+1}^T * \frac{\partial C}{\partial Z_{l+1}} \right)}_{\text{Looks forward}} * \underbrace{\sigma'(Z_l) * A_{l-1}^T}_{\text{Looks backward}}$$

5. **Update weights according to gradient descent**

$$W_{new} = W - \eta \nabla C$$

Improving speed at calculating ∇C

- Need to compute the gradient for each training example x and then average

$$- C = \frac{1}{n} \sum_x C_x \rightarrow \nabla C = \frac{1}{n} \sum_x \nabla C_x$$


- With typical $n > 10^4$ this is extremely slow

→ Stochastic gradient descent

- Randomly choose sample $m \sim 10^2$ of X (mini-batch)
- $\nabla C \approx \frac{1}{m} \sum_{j=1}^m \nabla C_{x_j}$
- Train on this mini-batch, then choose a new batch until having trained with all (1 epoch)
- Then start over and redo until training is finished (100-1000 epochs)
- **Less straight path, but still can be more time-efficient!**

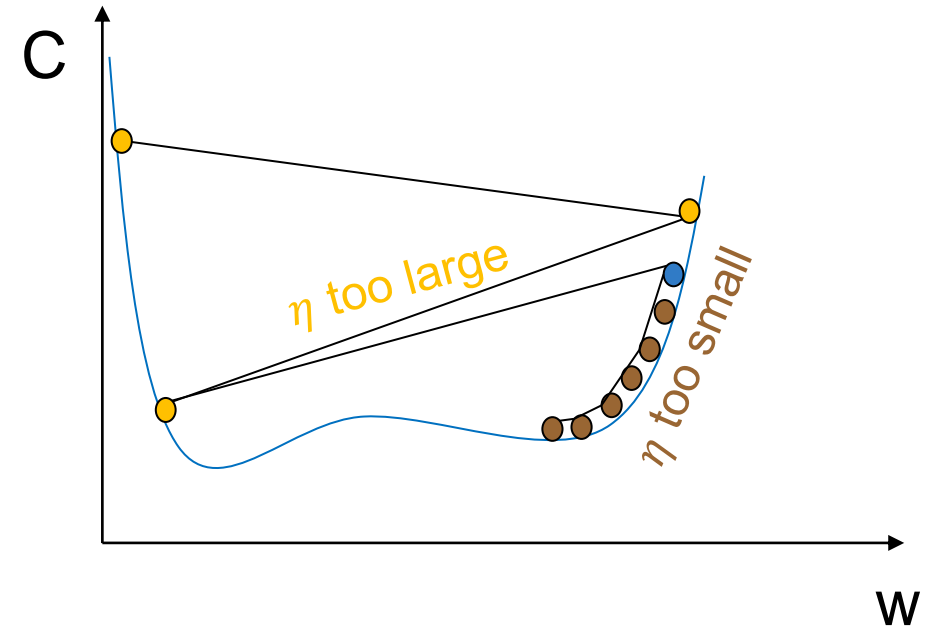
X : Complete data set ($n \sim 10^5$)

m	m	m	m	m	m
m	m	m	m	m	m
m	m	m	m	m	m
m	m	m	m	m	m
m	m	m	m	m	m



The learning rate

- Too small learning rate \rightarrow can get caught in local minima
- Too large learning rate \rightarrow Solution can diverge!
- Best solution: Adaptive learning rate based on
 - How large gradient is
 - How fast learning happens
 - Size of some weights
 - Momentum
 - ...



Energy use of training by backpropagation

- On average 10 ms time for one layer operation, t_{op} (forward + backwards pass)
- Hardware NVIDIA Tesla V100 GPU: 250W
- Ex: $L = 100$, $m = 100$, $N = 100\,000$, 100 epochs $\rightarrow t_{exec} = t_{op} L \frac{N}{m} e = 100\,000\,s \sim 27h$.
 $\rightarrow 25\text{ MJ energy consumption! } (0.25\text{kW} * 27h = 7\text{ kWh})$



1kWh



100 days of LED light, (2.5W)



~300 full charges of your cell phone


Vanishing gradient problem

- In back-propagation:

$$\frac{\partial C}{\partial W_l} = \frac{\partial C}{\partial Z_l} \frac{\partial Z_l}{\partial W_l} = \left(W_{l+1}^T * \frac{\partial C}{\partial Z_{l+1}} \right) * \underbrace{\sigma'(Z_l)}_{< 1} * \underbrace{A_{l-1}^T}_{< 1}$$

i.e. multiplication of numbers smaller than 1

→ Gradients vanish in deep networks






Deep Belief Nets

- A deep network that still avoids the vanishing gradient problem.
- NOT trained by backpropagation → unsupervised training
- Based on principles of Statistical Physics, Boltzmann statistics, probabilities, energies etc..
- Very brief introduction only...

Restricted Boltzmann Machines

- Two layer fully connected network with stochastic binary units
- Trained iteratively without labels by repeated (1) forward pass and (2) reconstruction pass
- Change weights as to minimize error between input and "reconstructed" input




Training Deep Belief Net

- Each bilayer RBM is trained individually
- Each hidden layer of RBM is visible layer of the next
- Train each RBM separately in sequence until the whole network is trained.
 - Unsupervised learning! No labels needed
- Finally: Finetuning with small labeled dataset



DBM as Autoencoder

- Stacking a "mirrored" DBM on the output → generate new data!
- Trained in the same way
- Useful for tons of things:
 - Data denoising
 - Generate new similar images/video based on data set
 - Anomaly detection
 - DeepFake
 - Find similar documents (Document retrieval)



Summary

- Artificial neural networks – abstraction of the learning principles of the brain (matrix math)
- Free parameters: weights of connections and biases of "neurons"
- Optimize parameter values to learn things
 - Backpropagation
 - Gradient descent
- Unsupervised learning → Deep Belief Nets
 - Generation of new data (autoencoder)

Exercise: Make your own Quiz

- Come up with a good quiz question based on the topics of L1-L3
- You have until the end of the lecture → Send it to mattias.borg@eit.lth.se
- Quiz will be posted on Canvas for you to practise on..

What is the chance that you win on the lottery?

1. Chance? I always win
2. 1 in 100 000
3. As good as dying in a plane crash
4. I will win when pigs can fly

Lecture 1 – Introduction
Memory cell

Lecture 2 – Current memories
DRAM
Flash
SRAM

Lecture 3 – Machine Learning Fundamentals
Backpropagation
Gradient descent
Restricted Boltzmann Machines
Deep Belief Networks